

Model selection for time series forecasting

Vadim Strijov
with

Anastasia Motrenko, Aleksandr Aduenko, Alexey Goncharov, Radoslav Neychev, and Maria Vladimirova

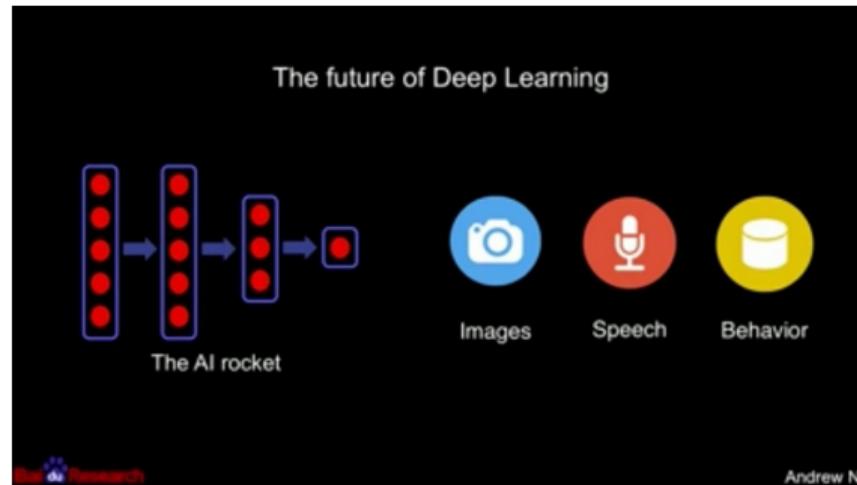
Moscow Institute of Physics and Technology

11th International Conference on Intelligent Data Processing: Theory and Applications
2016, October 10–14, Barcelona

Model selection for time series forecasting

The Internet of things is the world of networking devices (portables, vehicles, buildings) embedded with sensors and software.

- ▶ Environment and energy monitoring
- ▶ Medical and health monitoring
- ▶ Consumer support, sales monitoring
- ▶ Urban management and manufacturing



Case 1. Energy consumption and price forecasting, 1-day ahead hourly

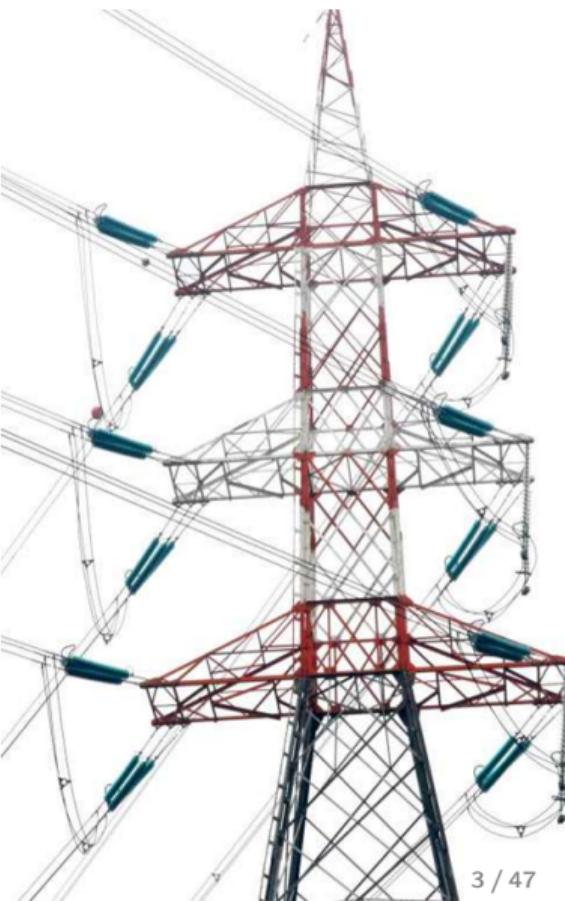
The components of multivariate time series with periodicity

Time series:

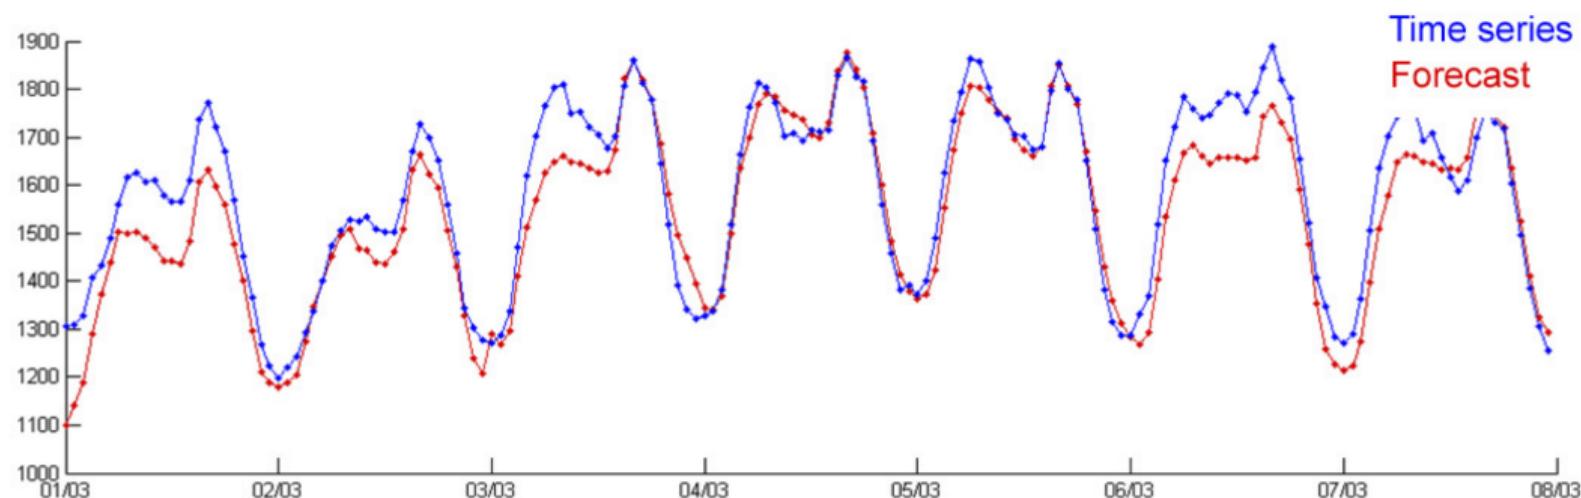
- ▶ energy price,
- ▶ consumption,
- ▶ daytime,
- ▶ temperature,
- ▶ humidity,
- ▶ wind force,
- ▶ holiday schedule.

Periodicity:

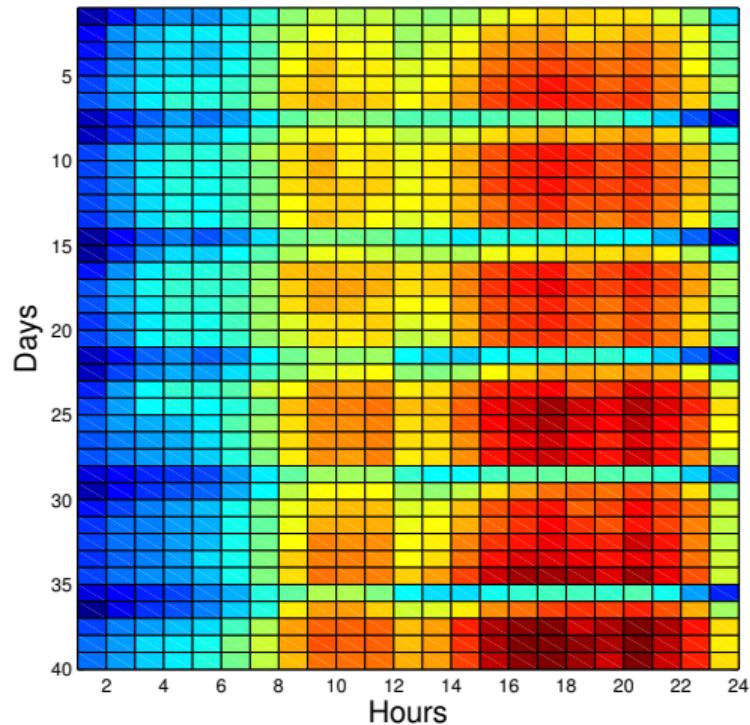
- ▶ one year seasons
(temperature,
daytime),
- ▶ one week,
- ▶ one day (working
day, week-end),
- ▶ a holiday,
- ▶ aperiodic events.



Energy consumption one-week forecast for each hour



The autoregressive matrix, five weeks



The autoregressive matrix and the linear model

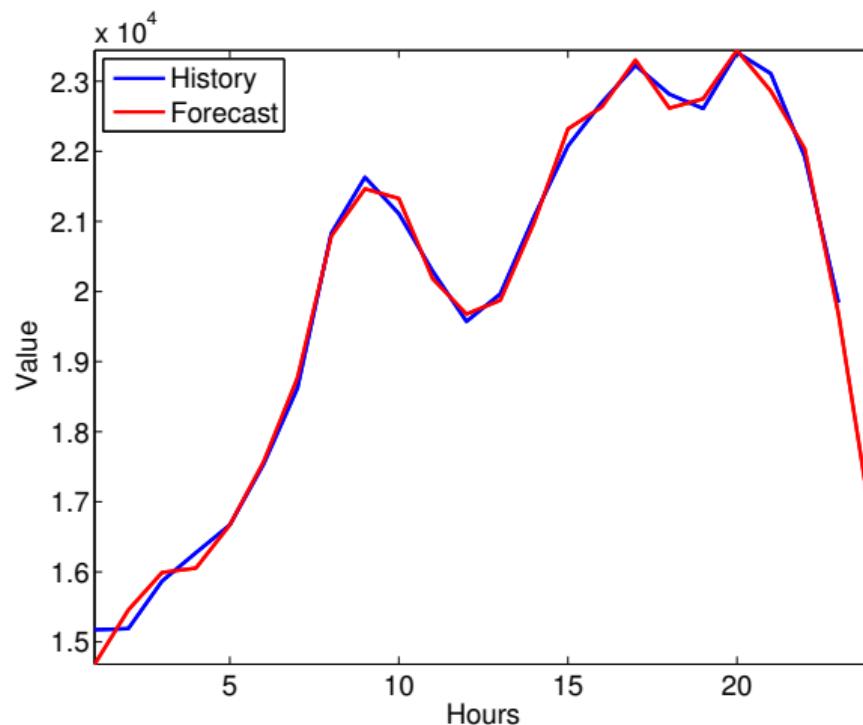
$$\mathbf{X}^*_{(m+1) \times (n+1)} = \left[\begin{array}{c|cccc} \hat{s}_T & s_{T-1} & \dots & s_{T-\kappa+1} \\ \hline s_{(m-1)\kappa} & s_{(m-1)\kappa-1} & \dots & s_{(m-2)\kappa+1} \\ \dots & \dots & \dots & \dots \\ s_{n\kappa} & s_{n\kappa-1} & \dots & s_{n(\kappa-1)+1} \\ \dots & \dots & \dots & \dots \\ s_\kappa & s_{\kappa-1} & \dots & s_1 \end{array} \right] = \left[\begin{array}{c|cc} \hat{s}_T & \mathbf{x}_{m+1} \\ \hline 1 \times 1 & 1 \times n \\ \mathbf{y} & \mathbf{X} \\ m \times 1 & m \times n \end{array} \right].$$

In terms of linear regression:

$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{X}, \mathbf{w}) = \mathbf{X}\mathbf{w},$$

$$\hat{y}_{m+1} = \hat{s}_T = \langle \mathbf{x}_{m+1}, \hat{\mathbf{w}} \rangle.$$

The one-day forecast: expected error is 3.1% working day, 3.7% week-end



The model $\hat{y} = f(\mathbf{X}, \mathbf{w})$ could be a linear model, neural network, deep NN, SVN, ...

The model performance criteria and forecast errors

Stability:

- ▶ the error does not change significantly following small changes in time series,
- ▶ the distribution of the model parameters does not change.

Complexity:

- ▶ the number of parameters (elements in superposition) is minimal,
- ▶ the minimum description length principle holds the William Ockham's rule.

Error: the residue $\varepsilon_j = \hat{y}_j - y_j$ for

- ▶ mean absolute error and (symmetric) mean absolute percent error

$$RSS = \sum_{j=1}^r \varepsilon_j^2, \quad MAPE = \frac{1}{r} \sum_{j=1}^r \frac{|\varepsilon_j|}{|y_j|}, \quad sMAPE = \frac{1}{r} \sum_{j=1}^r \frac{2|\varepsilon_j|}{|\hat{y}_j + y_j|}.$$

Case 2. Sales planning: to forecast the goods consumption

Retailers' daily routines:

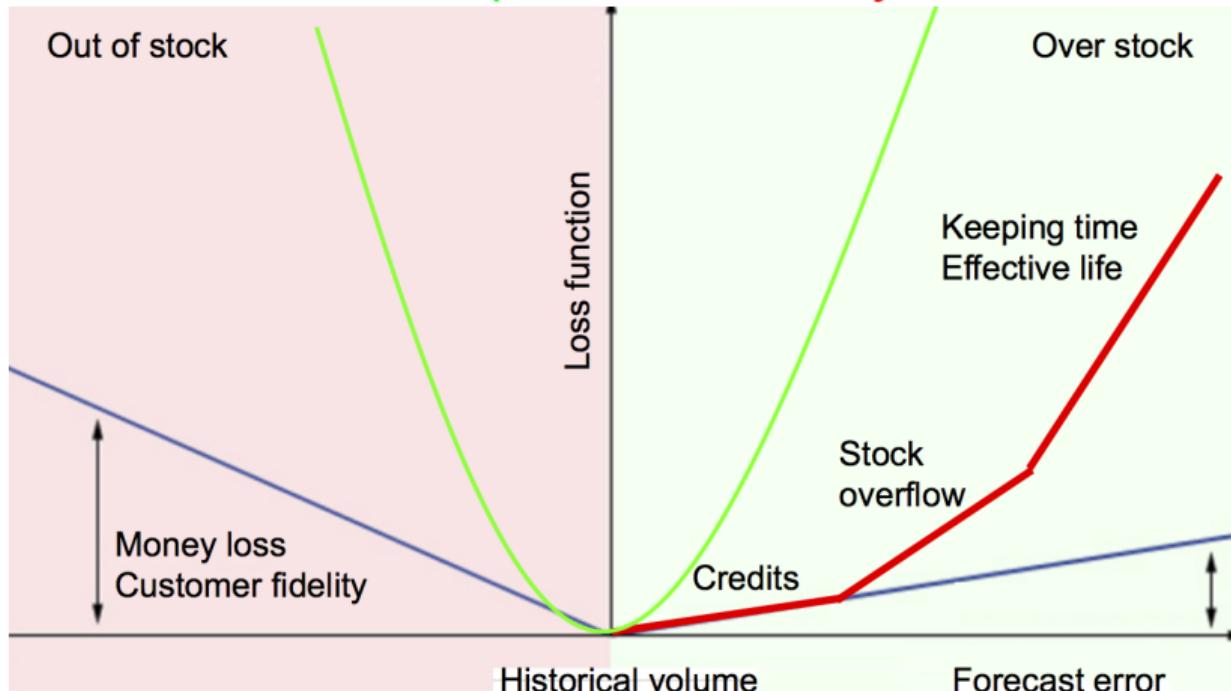
- ▶ custom inventory,
- ▶ calculation of optimal insurance stocks,
- ▶ consumer demand forecasting.



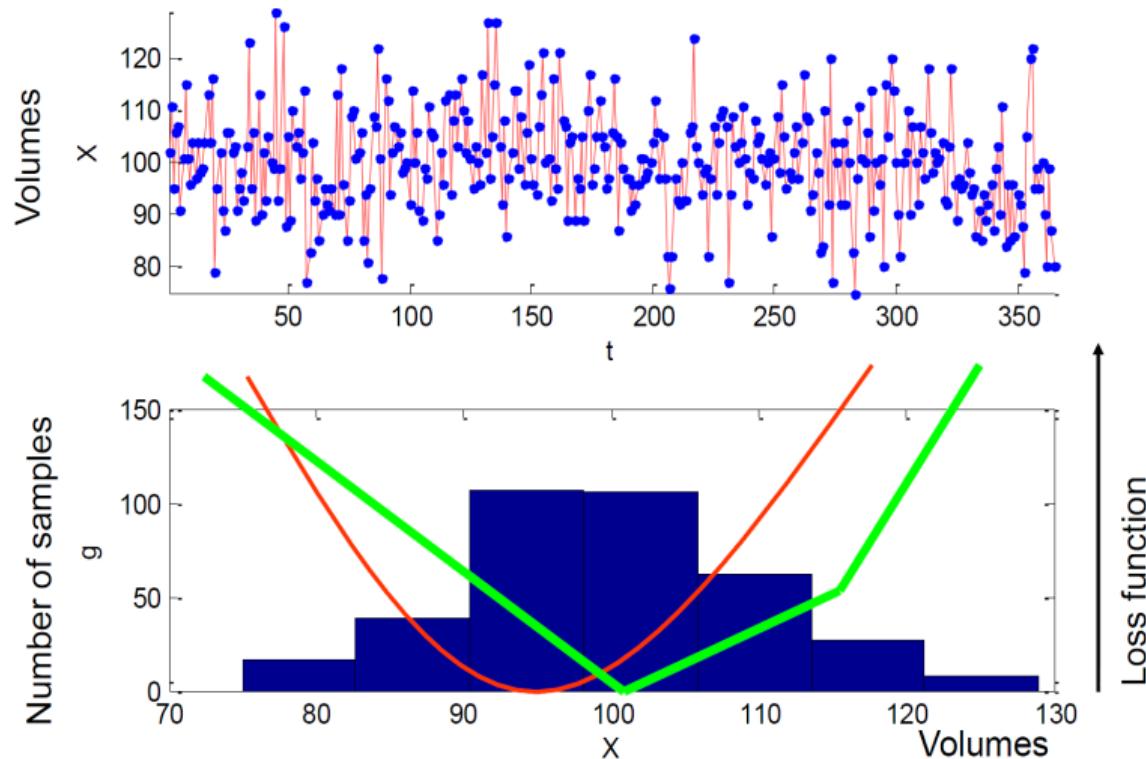
- ▶ There given historical time series of the volume off-takes: foodstuff.
- ▶ Let the time series be homoscedastic: its variance is time-constant.
- ▶ Minimizing the loss function one must forecast the next sample.

The performance criterion is **minimum loss of money**

Error functions: **quadratic**, **linear**, **asymmetric**.



The time series of residues and its histogram



There given historgam

$$H = \{X_i, g_i\}_{i=1}^m$$

and loss function

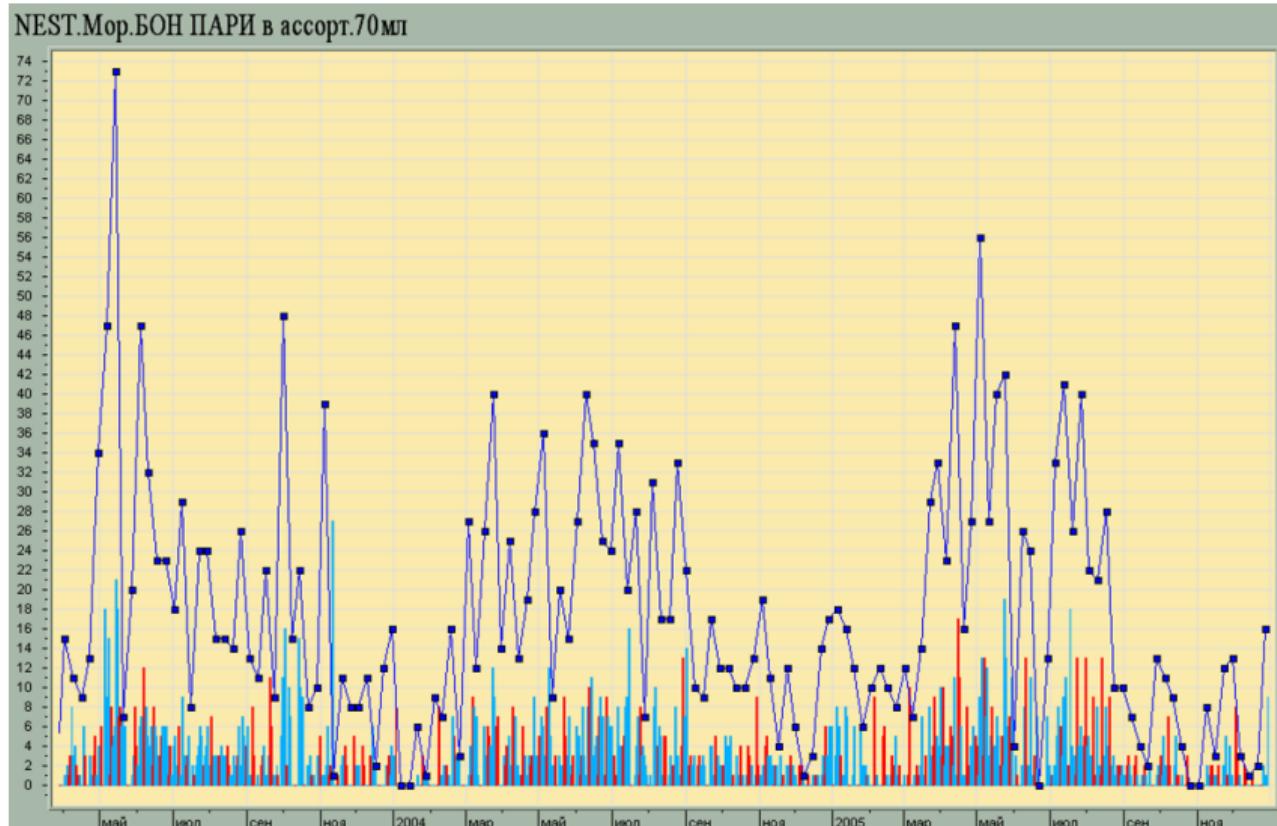
$$L = L(Z, X).$$

The optimal forecast is

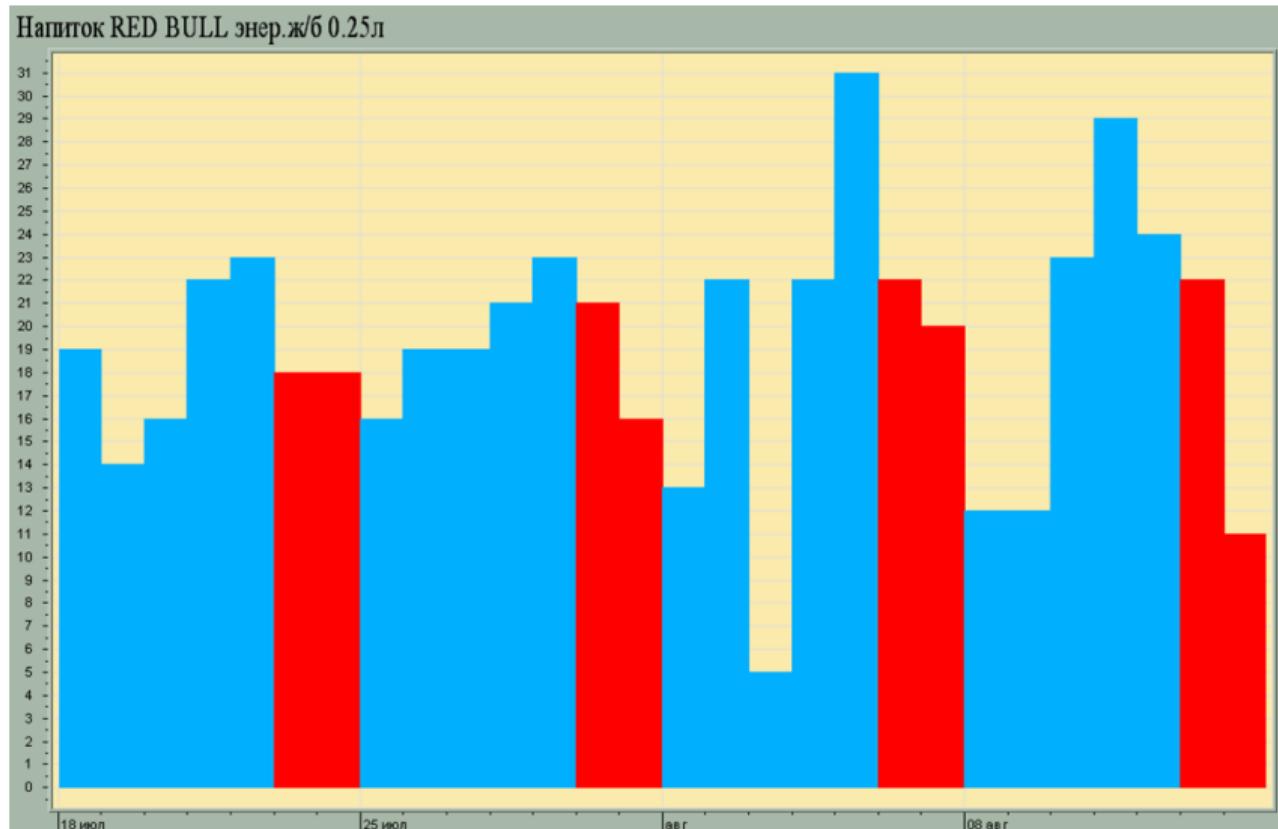
$$\tilde{X} = \arg \min_{Z \in \{X_1 \dots X_m\}} \sum_{i=1}^m g_i L(Z, X_i)$$

of this convolution.

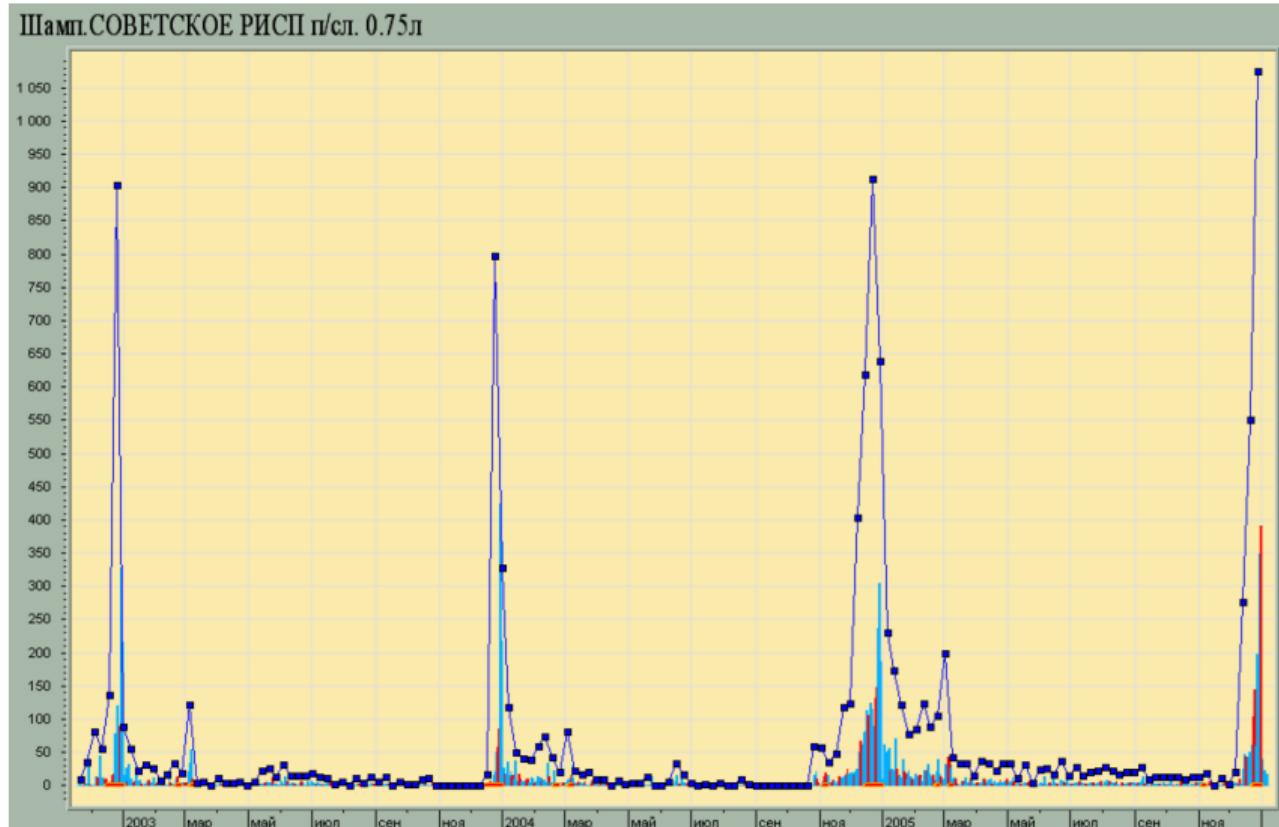
Candies: the seasonality and trend weekly over three years



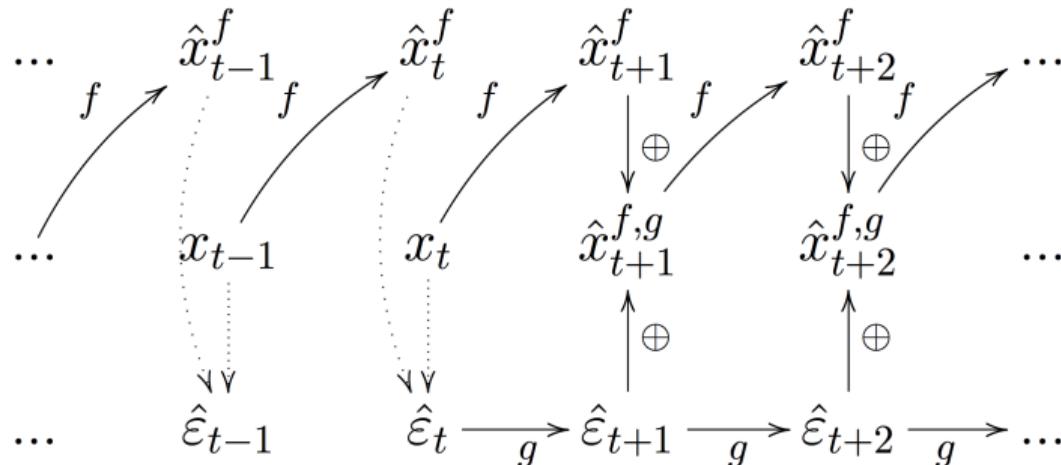
Beverage: the week periodicity daily over four weeks



Sparkling wine: holidays weekly over three years



Forecast the residues to boost performance



1. Model f forecasts $n(g)$ history ends $\hat{x}_t^f, \dots, \hat{x}_{t-n(g)+1}^f$ for one sample.
2. Compute $n(g)$ residues $\hat{\varepsilon}_t, \dots, \hat{\varepsilon}_{t-n(g)+1}$ as $\hat{\varepsilon}_{t-k} = x_{t-k} - \hat{x}_{t-k}^f$.
3. Function g forecasts residues $\hat{\varepsilon}_{t+i}$ ahead $\max(i)$ time-ticks.
4. Combine forecasts $\hat{x}_{t+i}^{f,g} = \hat{x}_{t+i}^f + \hat{\varepsilon}_{t+i}$ computing f for each sample \hat{x}_{t+i}^f .

Case 3. Forecasting volumes of Russian railways freight transportation

Keep a hierarchical structure of time series without loosing performance

Forecast with hierarchical aggregation of

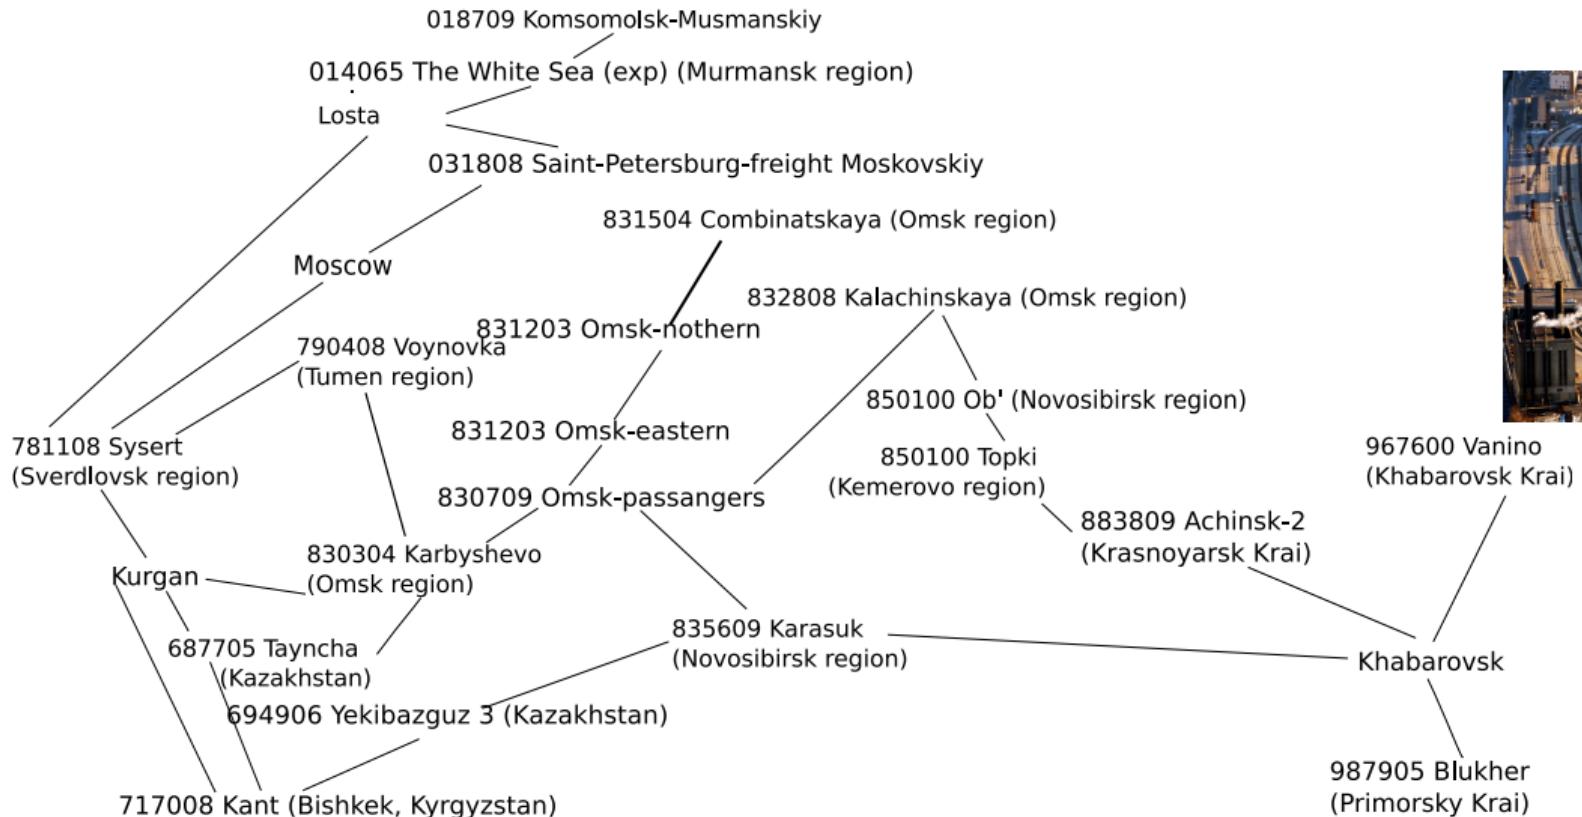
- ▶ types of freight in
- ▶ stations, regions, and roads,
- ▶ for a day, week, month, and quarter,
- ▶ counting all combinations above.

Satisfy the conditions:

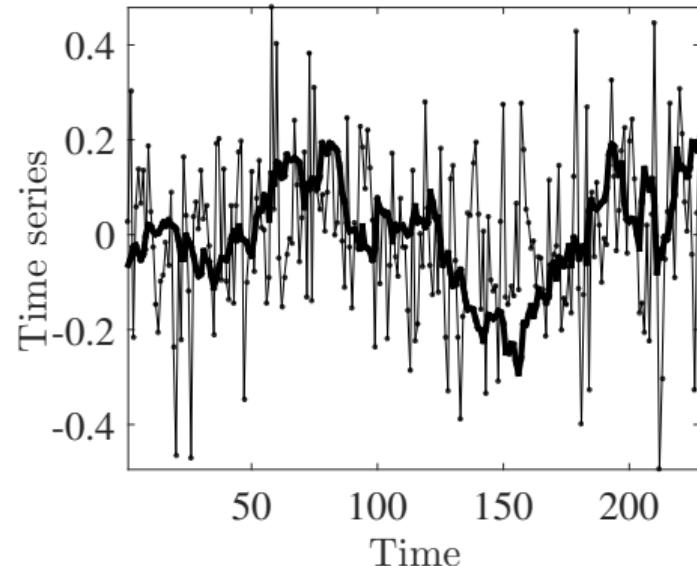
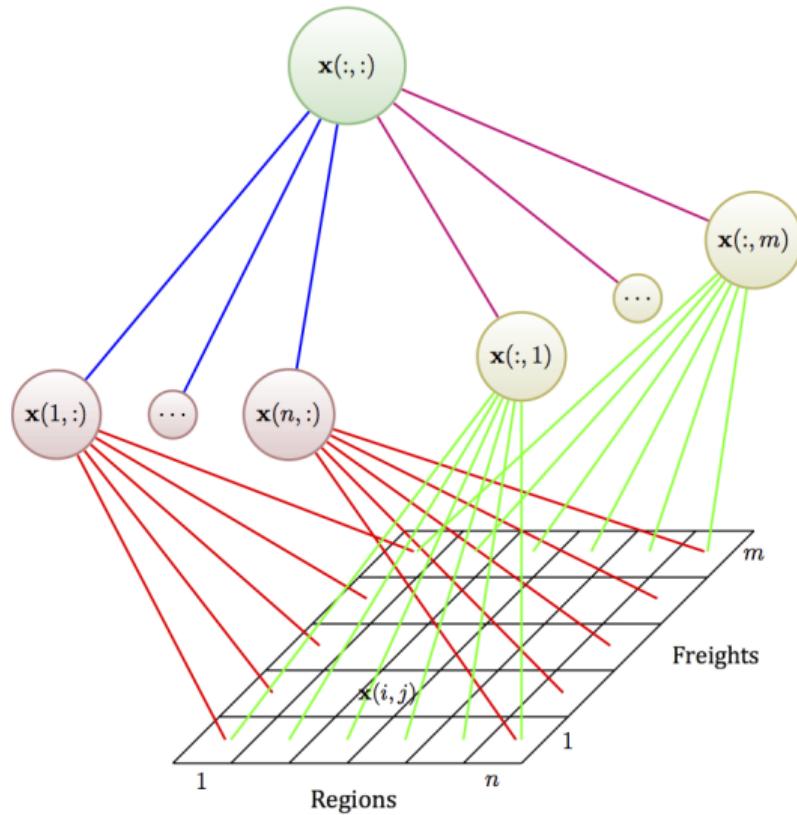
- ▶ minimize error,
- ▶ incorporate important external factors,
- ▶ respect hierarchical structure,
- ▶ do not exceed physical bounds of forecast values.



The railroad map counts ~ 78 regions, ~ 4000 stations, and ~ 100 rail-yards



Independent forecasts might be inconsistent with aggregated ones



Hierarchical data χ , independent forecasts $\hat{\chi}$ and reconciled forecasts $\hat{\varphi}$

$$\chi_t = \begin{pmatrix} x_t(:, :) \\ \dots \\ x_t(n, 1) \\ \dots \\ x_t(n, m) \end{pmatrix}, \quad \hat{\chi} = \begin{pmatrix} \hat{x}(:, :) \\ \dots \\ \hat{x}(n, 1) \\ \dots \\ \hat{x}(n, m) \end{pmatrix}, \quad \hat{\varphi} = \begin{pmatrix} \hat{y}(:, :) \\ \dots \\ \hat{y}(n, 1) \\ \dots \\ \hat{y}(n, m) \end{pmatrix}.$$

Consistency condition $S\chi_t = \mathbf{0}$, $t = 1, \dots, T$, where the link matrix S of size $(2 + n + m) \times (1 + n + m + nm)$ has the form

$$S = \left(\begin{array}{c|cccc|cccc|cccc|cccc|cccc} -1 & 1 & \dots & 1 & 0 & \dots & 0 & 0 & 0 & \dots & 0 & \dots & 0 & 0 & \dots & 0 \\ -1 & 0 & \dots & 0 & 1 & \dots & 1 & 0 & 0 & \dots & 0 & \dots & 0 & 0 & \dots & 0 \\ \hline 0 & -1 & \dots & 0 & 0 & \dots & 0 & 1 & 1 & \dots & 1 & \dots & 0 & 0 & \dots & 0 \\ \dots & \dots \\ 0 & 0 & \dots & -1 & 0 & \dots & 0 & 0 & 0 & \dots & 0 & \dots & 1 & 1 & \dots & 1 \\ \hline 0 & 0 & \dots & 0 & -1 & \dots & 0 & 1 & 0 & \dots & 0 & \dots & 1 & 0 & \dots & 0 \\ \dots & \dots \\ 0 & 0 & \dots & 0 & 0 & \dots & -1 & 0 & 0 & \dots & 1 & \dots & 0 & 0 & \dots & 1 \end{array} \right).$$

The regions

$$x_t(:, :) = \sum_{i=1}^n x_t(i, :);$$

The freights

$$x_t(:, :) = \sum_{j=1}^m x_t(:, j);$$

Freights given a region

$$x_t(i, :) = \sum_{j=1}^m x_t(i, j),$$

$$i = 1, \dots, n;$$

Regions given a freight

$$x_t(:, j) = \sum_{i=1}^n x_t(i, j),$$

$$j = 1, \dots, m;$$

$$t = 1, \dots, T.$$

Performance of independent and reconciliated forecasts

There given link matrix \mathbf{S} , admissible sets \mathbf{A}, \mathbf{B} and independent forecasts $\hat{\boldsymbol{\chi}} \notin \mathbf{A}, \hat{\boldsymbol{\chi}} \in \mathbf{B}$ to make reconciliated forecast $\hat{\boldsymbol{\varphi}}$ subject to

- consistency $\hat{\boldsymbol{\varphi}} \in \mathbf{A}, \mathbf{A} = \{\boldsymbol{\chi} \in \mathbb{R}^d \mid \mathbf{S}\boldsymbol{\chi} = \mathbf{0}\}$,
- physical limitations $\hat{\boldsymbol{\varphi}} \in \mathbf{B}$,
- precision $I_h(\boldsymbol{\chi}_{T+1}, \hat{\boldsymbol{\varphi}}) \leq I_h(\boldsymbol{\chi}_{T+1}, \hat{\boldsymbol{\chi}})$ for any hierarchical data $\boldsymbol{\chi}_{T+1} \in \mathbf{A} \cap \mathbf{B}$.

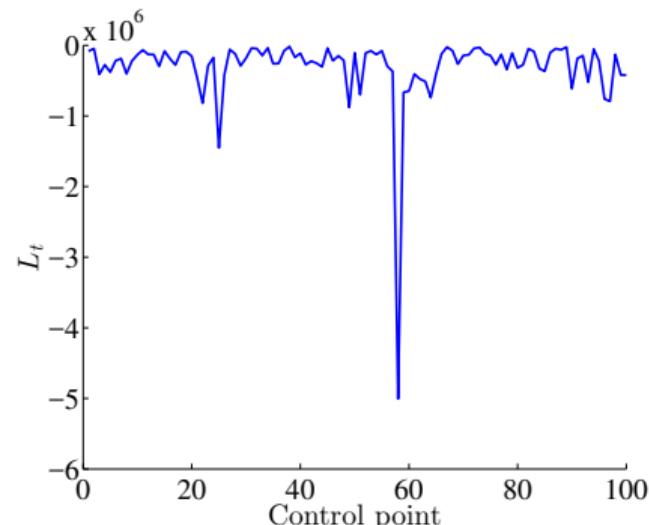
Theorem [Maria Stenina, 2014]

Given listed conditions, the projection vector

$$\hat{\boldsymbol{\varphi}} = \boldsymbol{\chi}_{proj} = \arg \min_{\boldsymbol{\chi} \in \mathbf{A} \cap \mathbf{B}} I_h(\boldsymbol{\chi}, \hat{\boldsymbol{\chi}})$$

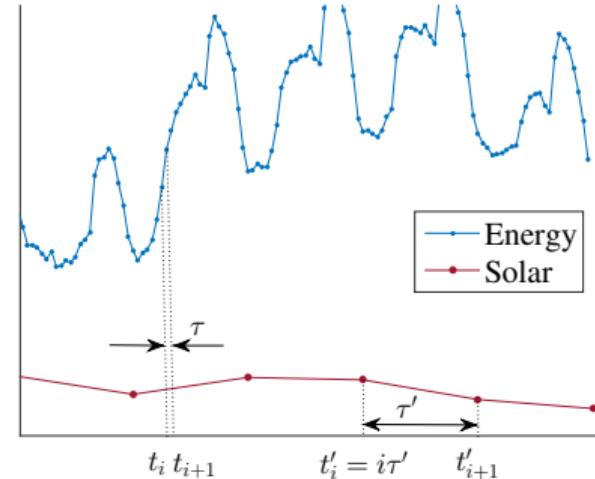
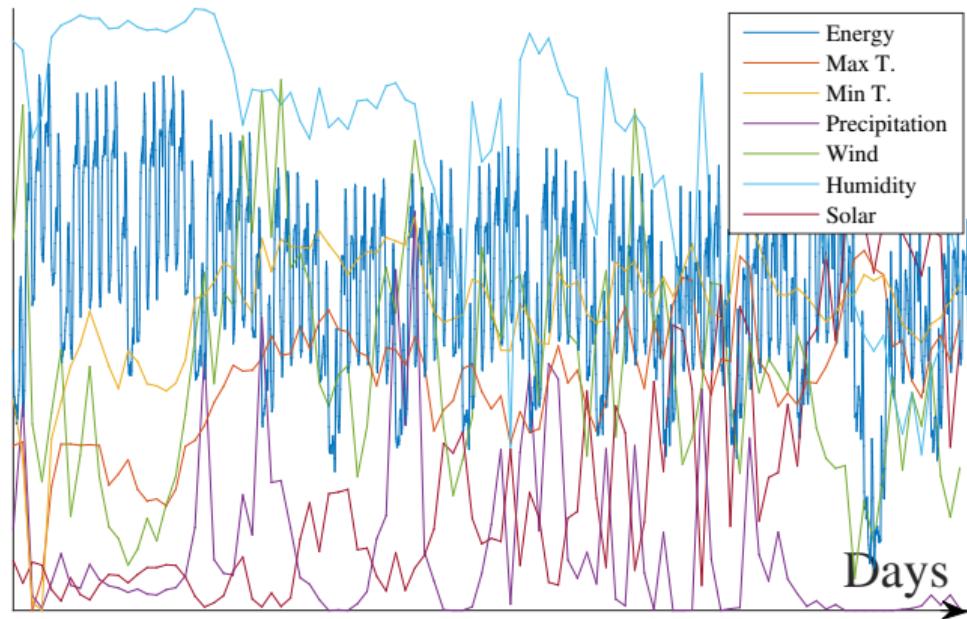
is guaranteed to satisfy the requirements of consistency, physical limitations and precision.

The solution of the optimization problem $\hat{\boldsymbol{\varphi}} = \arg \min_{\boldsymbol{\chi} \in \mathbf{A} \cap \mathbf{B}} \|\boldsymbol{\chi} - \hat{\boldsymbol{\chi}}\|_2^2$ demonstrates decrease in loss for all control samples.



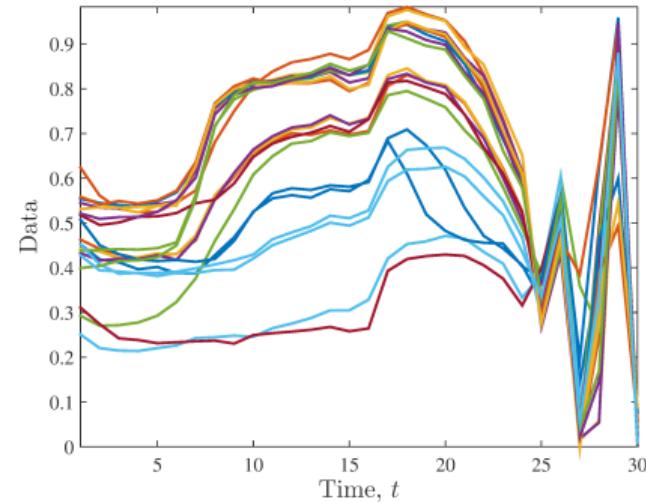
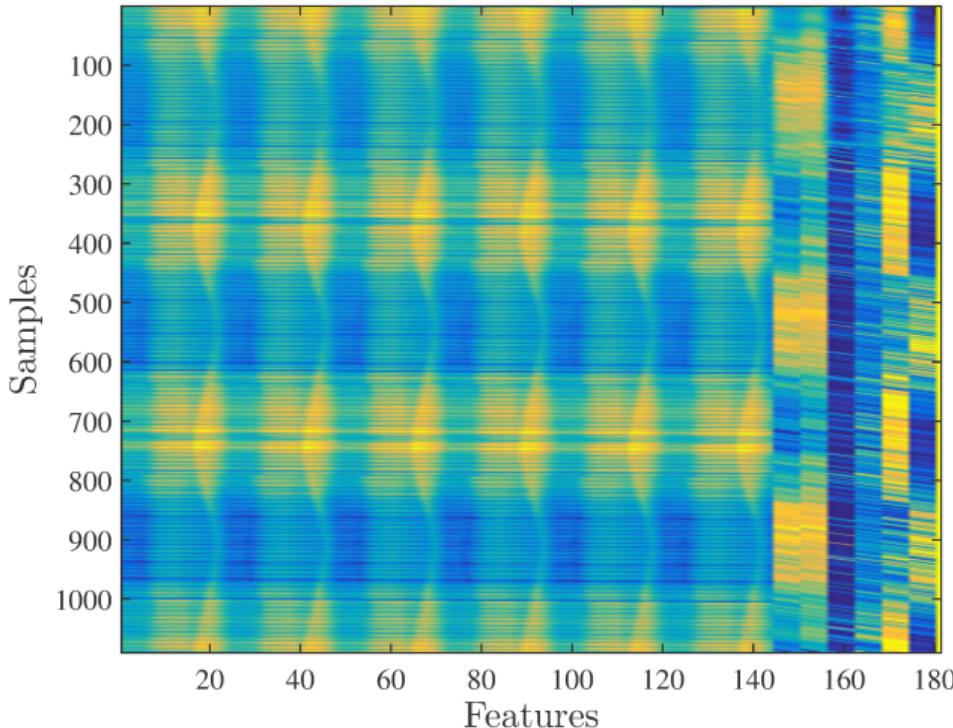
$$L_t = \|\boldsymbol{\chi}_t - \hat{\boldsymbol{\varphi}}\|_2^2 - \|\boldsymbol{\chi}_t - \hat{\boldsymbol{\chi}}\|_2^2$$

Case 4. Internet of things, multiscale dataset for vector forecasting



Each real-valued time series $\mathbf{s} = [s_1, \dots, s_i, \dots, s_T]$, $s_i = s(t_i)$, $0 \leq t_i \leq t_{\max}$
is a sequence of observations of some real-valued signal $s(t)$ with its own sampling rate τ .

To boost the forecast quality include external variables



Generate features with nonparametric transformation functions

Univariate

Formula	Output dimension
\sqrt{x}	1
$x\sqrt{x}$	1
$\arctan x$	1
$\ln x$	1
$x \ln x$	1

Bivariate

Plus	$x_1 + x_2$
Minus	$x_1 - x_2$
Product	$x_1 \cdot x_2$
Division	$\frac{x_1}{x_2}$
	$x_1 \sqrt{x_2}$
	$x_1 \ln x_2$

Nonparametric aggregation: sample statistics

Nonparametric transformations include basic data statistics:

- ▶ Sum or average value of each row \mathbf{x}_i , $i = 1, \dots, m$:

$$\phi_i = \sum_{j=1}^n x_{ij}, \text{ or } \phi'_i = \frac{1}{n} \sum_{j=1}^n x_{ij}.$$

- ▶ Min and max values: $\phi_i = \min_j x_{ij}$, $\phi'_i = \max_j x_{ij}$.
- ▶ Standard deviation:

$$\phi_i = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (x_{ij} - \text{mean}(\mathbf{x}_i))^2}.$$

- ▶ Data quantiles: $\phi_i = [X_1, \dots, X_K]$, where

$$\sum_{j=1}^n [X_{k-1} < x_{ij} \leq X_k] = \frac{1}{K}, \text{ for } k = 1, \dots, K.$$

Nonparametric transformations: Haar's transform

Applying Haar's transform produces multiscale representations of the same data.

Assume that $n = 2^K$ and init $\phi_{i,j}^{(0)} = \phi'_{i,j}^{(0)} = x_{ij}$ for $j = 1, \dots, n$.

To obtain coarse-graining and fine-graining of the input feature vector \mathbf{x}_i , for $k = 1, \dots, K$ repeat:

- ▶ data averaging step

$$\phi_{i,j}^{(k)} = \frac{\phi_{i,2j-1}^{(k-1)} + \phi_{i,2j}^{(k-1)}}{2}, \quad j = 1, \dots, \frac{n}{2^k},$$

- ▶ and data differencing step

$$\phi'_{i,j}^{(k)} = \frac{\phi'_{i,2j}^{(k-1)} - \phi'_{i,2j-1}^{(k-1)}}{2}, \quad j = 1, \dots, \frac{n}{2^k}.$$

The resulting multiscale feature vectors are $\phi_i = [\phi_i^{(1)}, \dots, \phi_i^{(K)}]$ and $\phi'_i = [\phi'_i^{(1)}, \dots, \phi'_i^{(K)}]$.

Collection of parametric transformation functions

Function name	Formula	Output dim.	Num. of args	Num. of pars
Add constant	$x + w$	1	1	1
Quadratic	$w_2x^2 + w_1x + w_0$	1	1	3
Cubic	$w_3x^3 + w_2x^2 + w_1x + w_0$	1	1	4
Logarithmic sigmoid	$1/(w_0 + \exp(-w_1x))$	1	1	2
Exponent	$\exp x$	1	1	0
Normal	$\frac{1}{w_1\sqrt{2\pi}} \exp\left(\frac{(x-w_2)^2}{2w_1^2}\right)$	1	1	2
Multiply by constant	$x \cdot w$	1	1	1
Monomial	$w_1x^{w_2}$	1	1	2
Weibull-2	$w_1 w_2 x^{w_2-1} \exp -w_1 x^{w_2}$	1	1	2
Weibull-3	$w_1 w_2 x^{w_2-1} \exp -w_1(x - w_3)^{w_2}$	1	1	3
...

Models and features

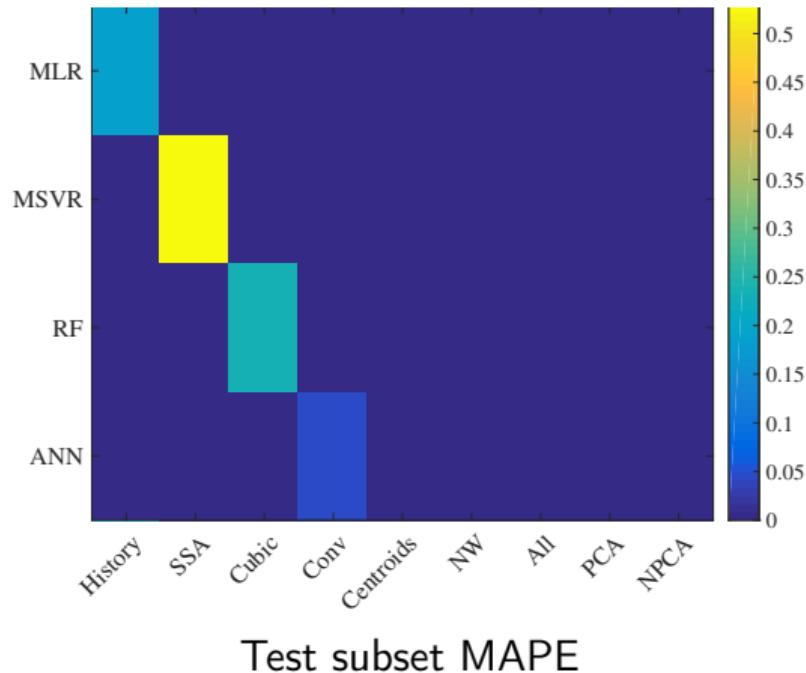
Models:

- ▶ Baseline method: $\hat{s}_i = s_{i-1}$.
- ▶ Multivariate linear regression (MLR) with l_2 -regularization. Regularization coefficient: 2.
- ▶ SVR with multiple output. Kernel type: RBF, $p_1: 2$, $p_2: 0$, $\gamma: 0.5$, $\lambda: 4$.
- ▶ Feed-forward ANN with single hidden layer, size: 25
- ▶ Random forest (RF). Number of trees: 25 , number of variables for each decision split: 48.

Feature combinations:

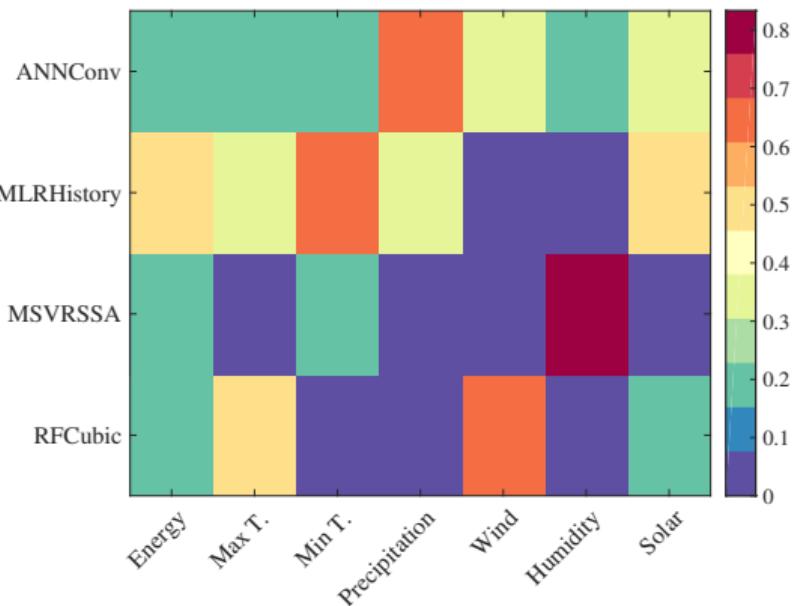
- ▶ History: the standard regression-based forecast with no additional features.
- ▶ SSA, Cubic, Conv, Centroids, NW: history + a particular feature.
- ▶ All: all of the above, with no feature selection.
- ▶ PCA and NPCA: all generation strategies with feature selection.

Feature analysis

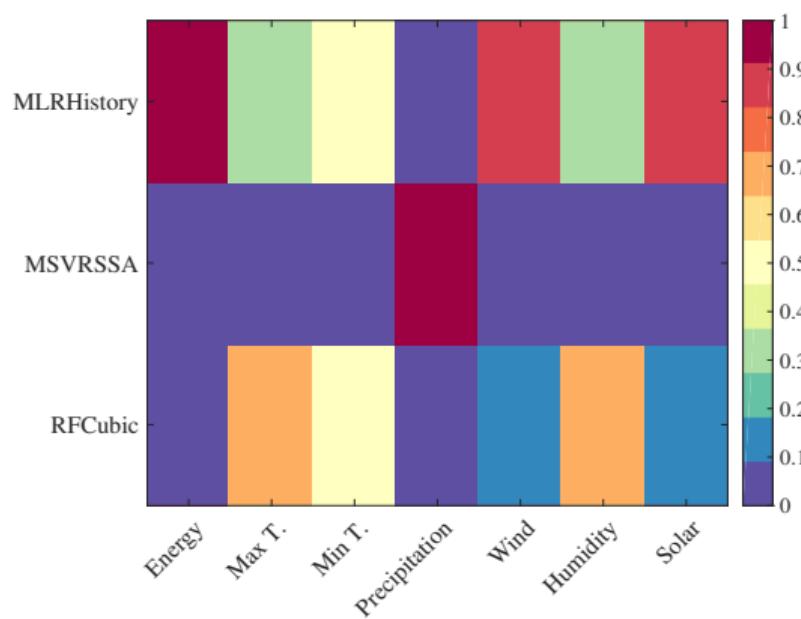


Ratio of times each combination of model and feature performed best for at least one of the time series (7) or error functions (6), all (6) data sets ($6 \times 7 \times 6 = 252$ cases).

Best models



Test subset residues



Standard deviations of test subset residues

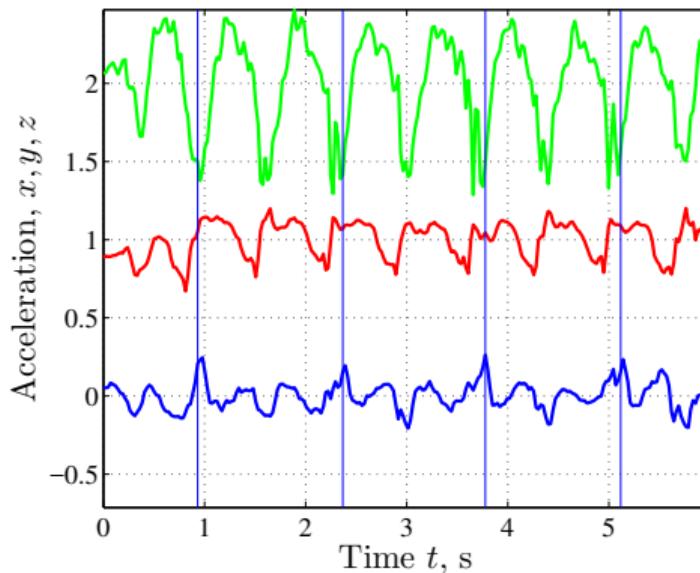
Case 5. Classification of human physical activity with mobile devices



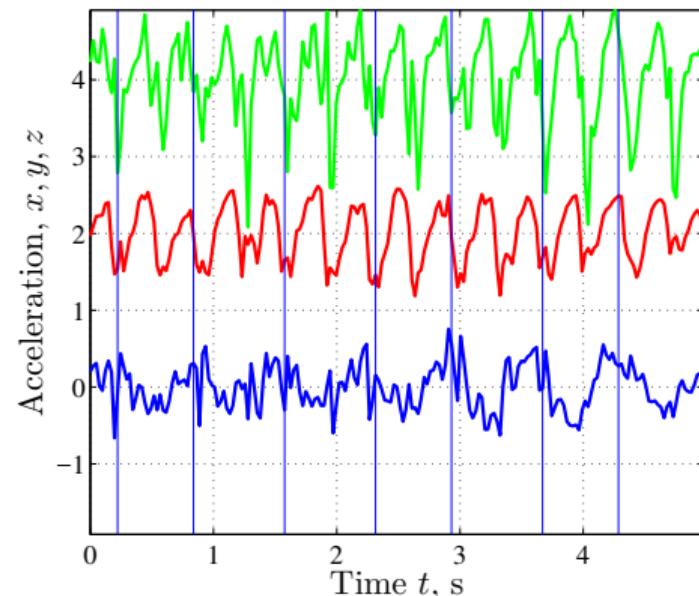
3D-projection of acceleration time series to spatial axis

$$\mathbf{x} = \{acc_x(t); acc_y(t); acc_z(t)\}_{t=1}^n \mapsto \mathbf{y} \in \mathbb{R}^S.$$

Slow walking



Jogging



Local transformations for deep learning neural network

Model $\mathbf{f} = \mathbf{a}(\mathbf{h}_N(\dots \mathbf{h}_1(\mathbf{x})))(\mathbf{w})$ contains autoencoders \mathbf{h}_k and softmax classifier \mathbf{a} :

$$\mathbf{f}(\mathbf{w}, \mathbf{x}) = \frac{\exp(\mathbf{a}(\mathbf{x}))}{\sum_j \exp(a_j(\mathbf{x}))}, \quad \mathbf{a}(\mathbf{x}) = \mathbf{W}_2^\top \tanh(\mathbf{W}_1^\top \mathbf{x}), \quad \mathbf{h}_k(\mathbf{x}) = \sigma(\mathbf{W}_k \mathbf{x} + \mathbf{b}_k),$$

where \mathbf{w} minimizes the error function.

Feature generation by local transformations:

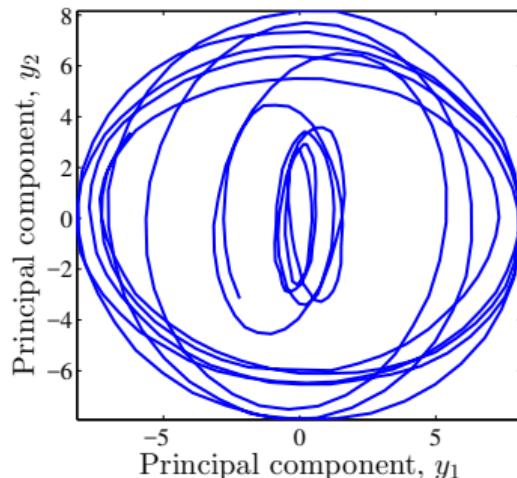
- ▶ parameters of SSA approximation of the time series \mathbf{x} ,
- ▶ FFT of \mathbf{x} ,
- ▶ parameters of polynomial/spline approximation,

could reduce complexity this model down to complexity of logistic regression.

Parameters of the local models: SSA

For time series \mathbf{s} construct the Hankel matrix with a period k and shift p , so that for $\mathbf{s} = [s_1, \dots, s_T]$ the matrix

$$\mathbf{H}^* = \begin{bmatrix} s_T & \dots & s_{T-k+1} \\ \vdots & \ddots & \vdots \\ s_{k+2} & \dots & s_2 \\ s_k & \dots & s_1 \end{bmatrix}.$$



Reconstruct the regression to the first column of the matrix $\mathbf{H}^* = [\mathbf{h}, \mathbf{H}]$ and denote its least square parameters as the feature vector

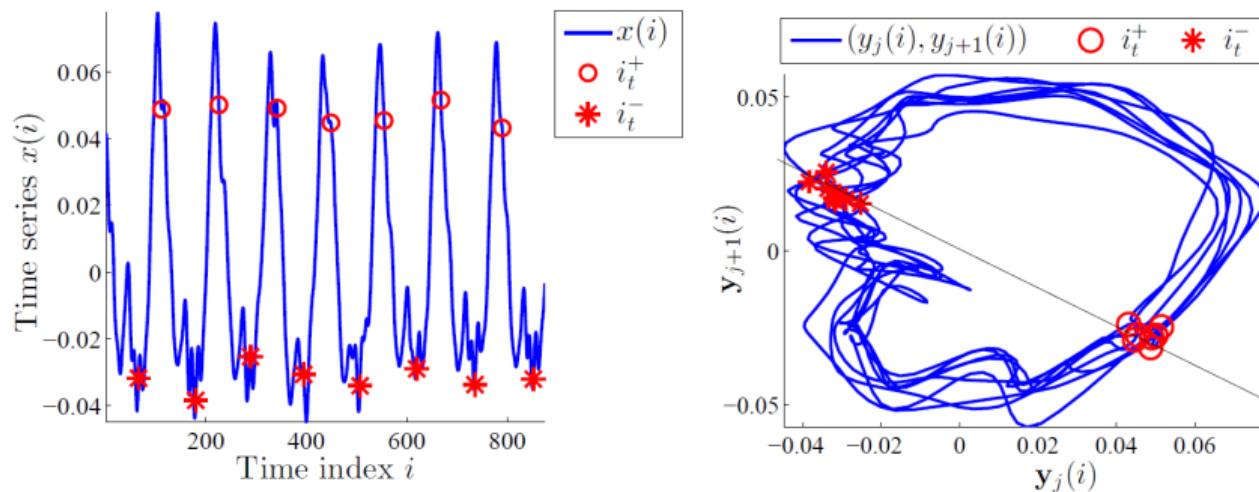
$$\phi(\mathbf{s}) = \arg \min \|\mathbf{h} - \mathbf{H}\phi\|_2^2.$$

For the original feature vector $\mathbf{x} = [\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(Q)}]$ use the parameters $\phi(\mathbf{x})$ as an item.

Human gate detection with time series segmentation

Find dissection of the trajectory of principal components $\mathbf{y}_j = \mathbf{H}\mathbf{v}_j$, where \mathbf{H} is the Hankel matrix and \mathbf{v}_j are its eigenvectors:

$$\frac{1}{N} \mathbf{H}^T \mathbf{H} = \mathbf{V} \Lambda \mathbf{V}^T, \quad \Lambda = \text{diag}(\lambda_1, \dots, \lambda_N).$$



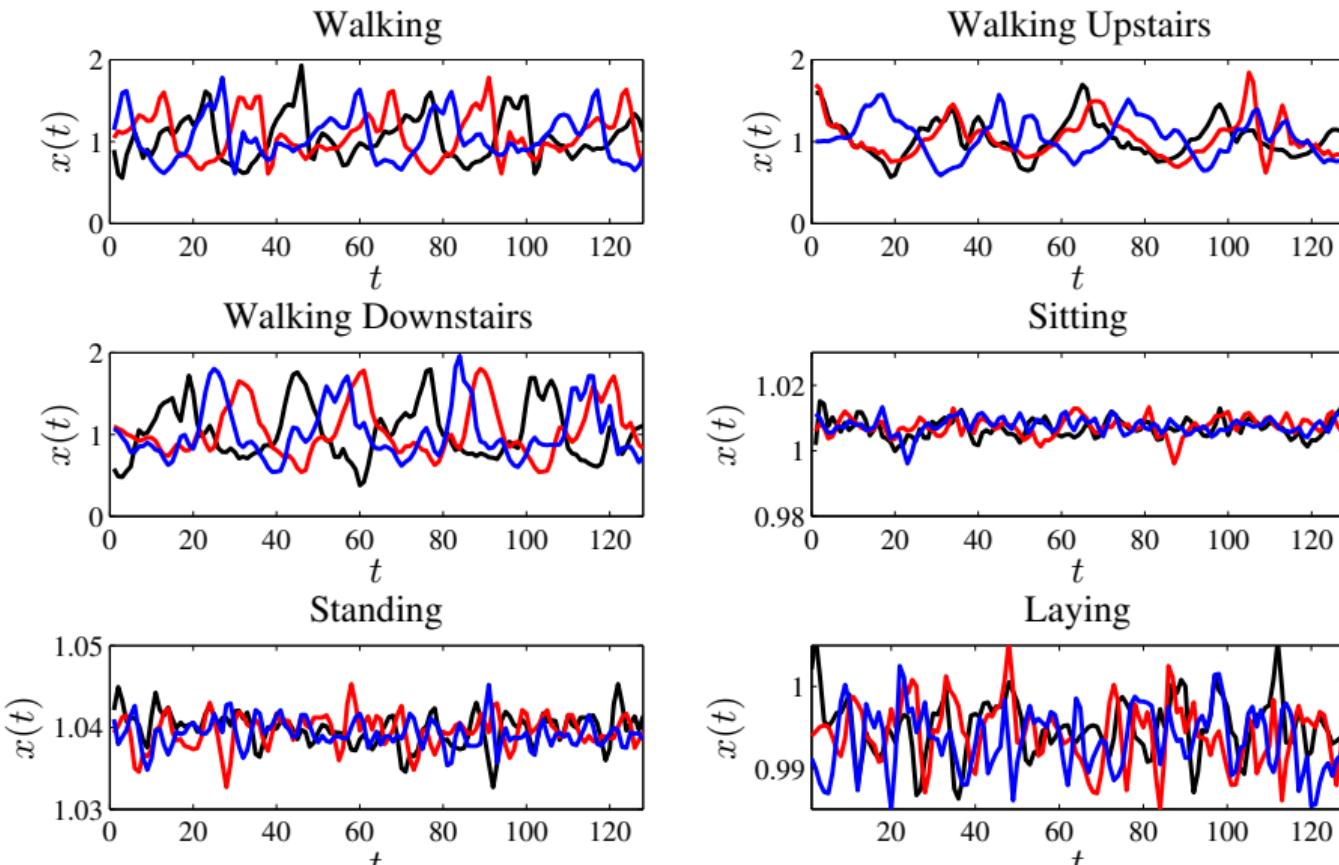
Metric features: distances to the centroids of local clusters

Apply kernel trick to the time series.

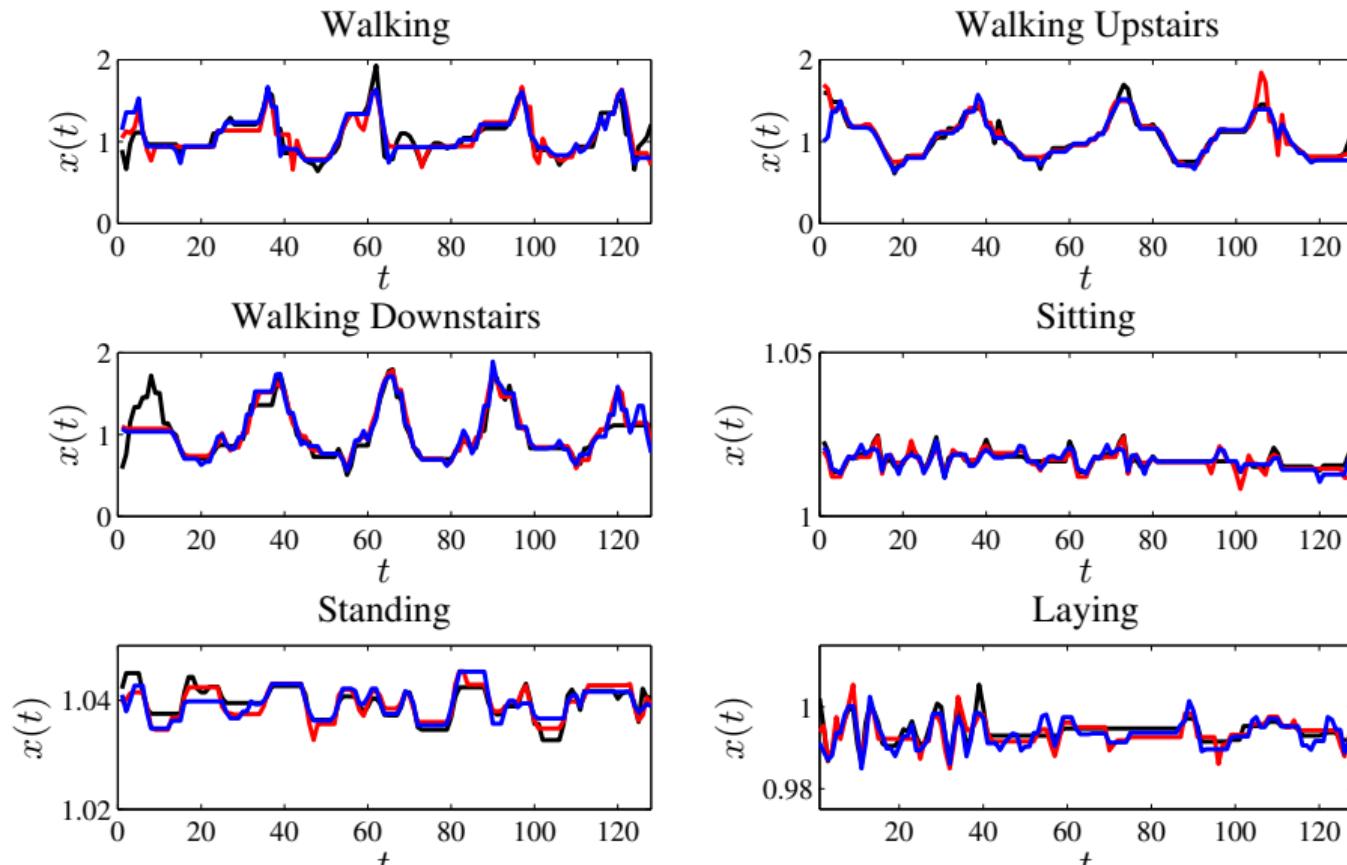
1. For objects \mathbf{x}_i from \mathbf{X} compute k -mean centroids \mathbf{c} .
2. Use distance function ρ to combine feature vector

$$\phi_i = [\rho(\mathbf{c}_1, \mathbf{x}_i), \dots, \rho(\mathbf{c}_p, \mathbf{x}_i)] \in \mathbb{R}_+^p.$$

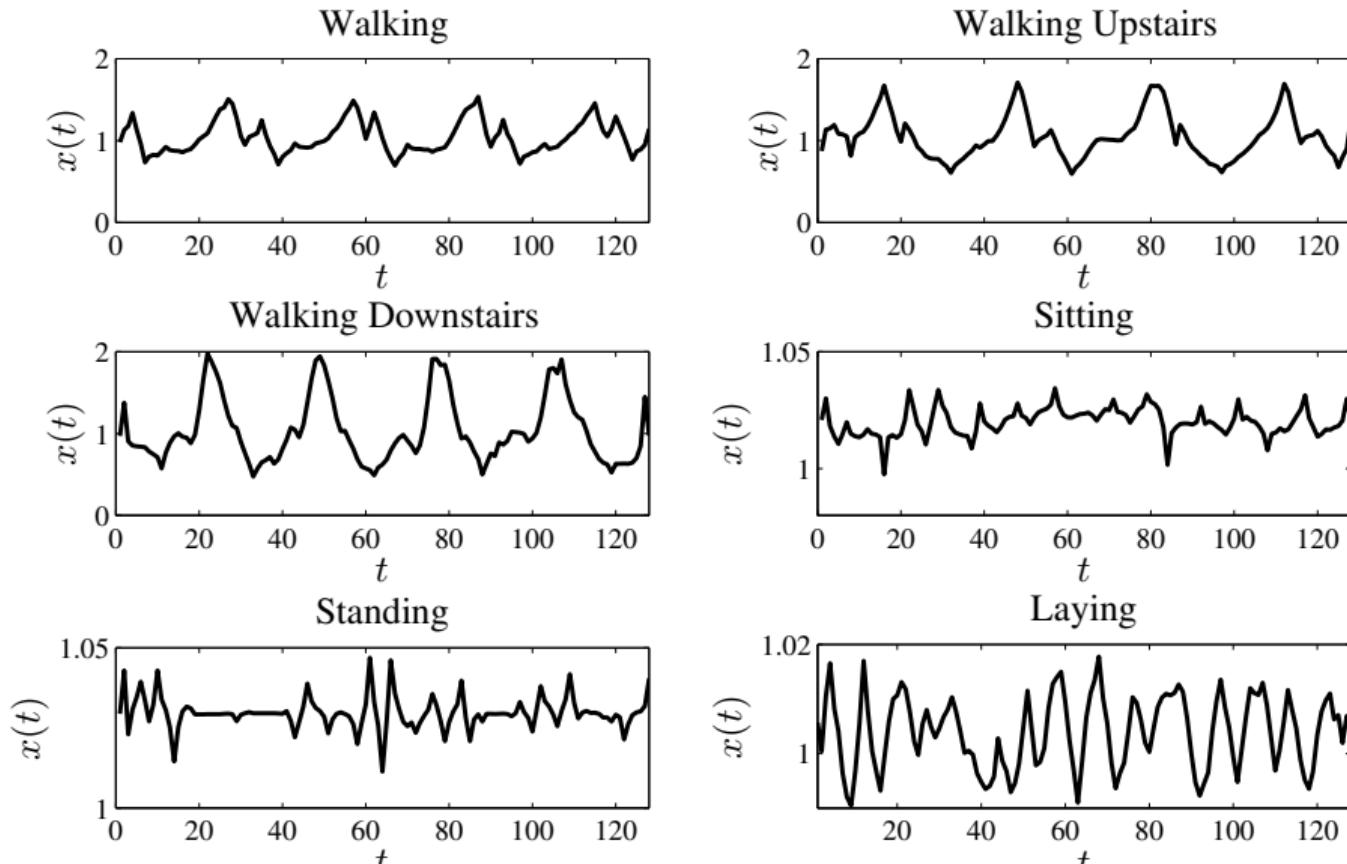
Computing distances to centroids of the time series: **sources**



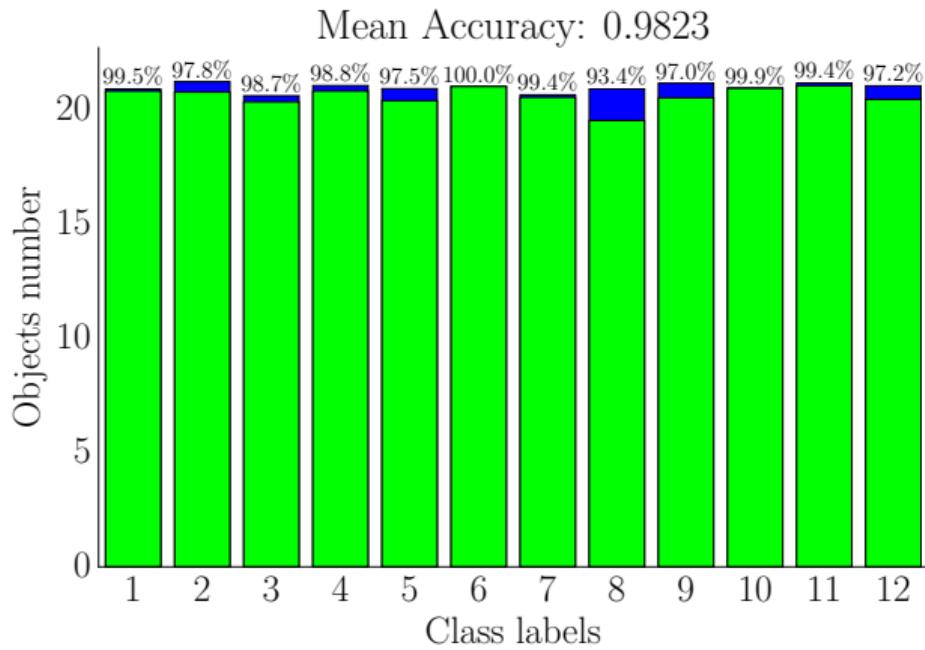
Computing distances to centroids of the time series: **aligned series**



Computing distances to centroids of the time series: **centroids**

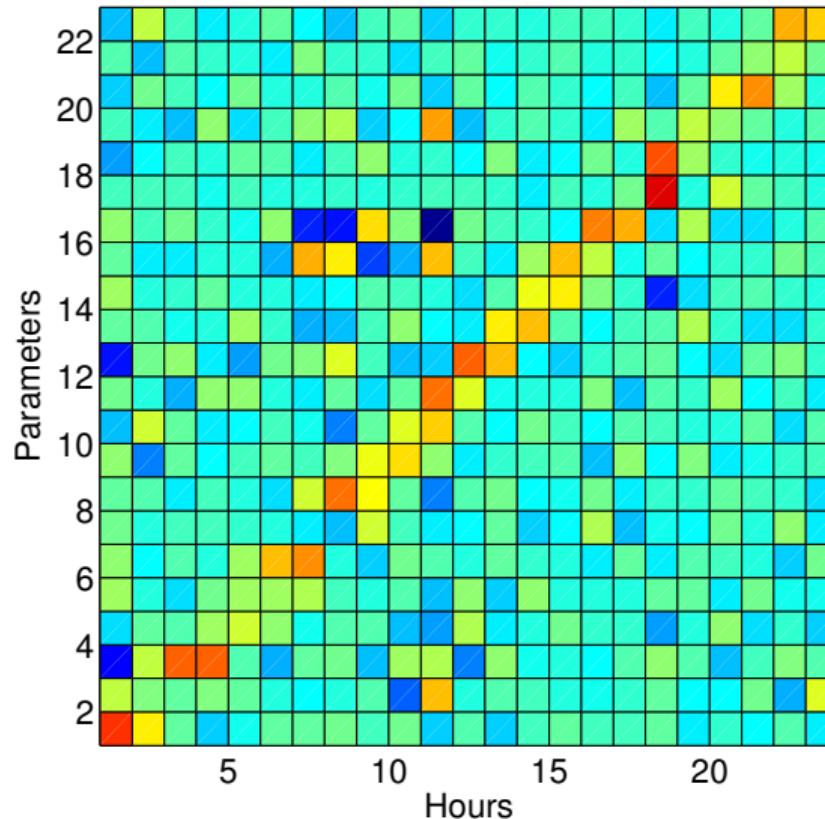


Performance of the human physical activities classification model



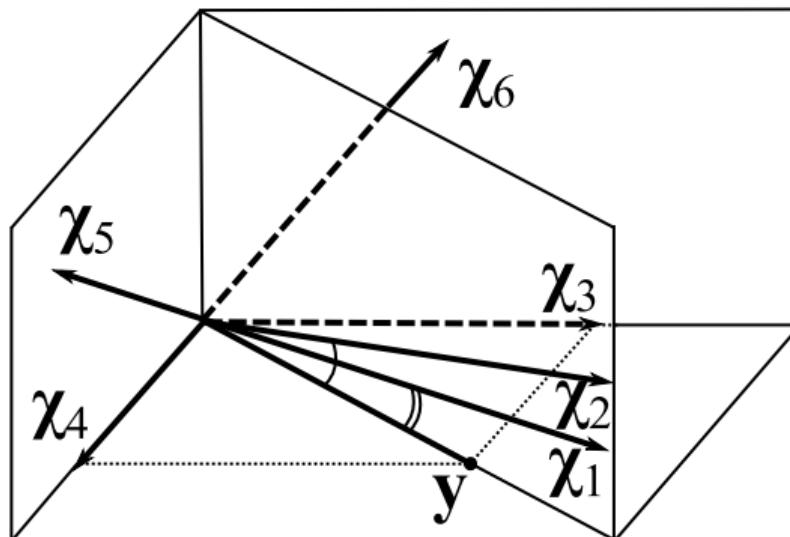
- 1) walk forward
- 2) walk left
- 3) walk right
- 4) go upstairs
- 5) go downstairs
- 6) run forward
- 7) jump up and down
- 8) sit and fidget
- 9) stand
- 10) sleep
- 11) elevator up
- 12) elevator down

Case 6. How many parameters must be used? Relations of 24 hourly models



Selection of a stable set of features of restricted size

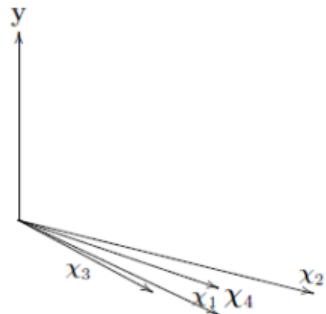
The sample contains multicollinear χ_1, χ_2 and noisy χ_5, χ_6 features, columns of the design matrix \mathbf{X} . We want to select two features from six.



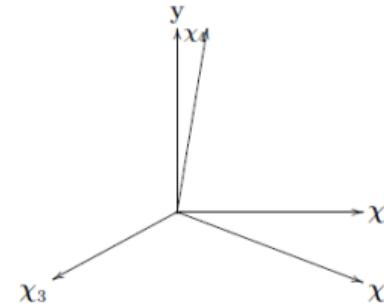
Stability and accuracy for a fixed complexity

The solution: χ_3, χ_4 is an orthogonal set of features minimizing the error function.

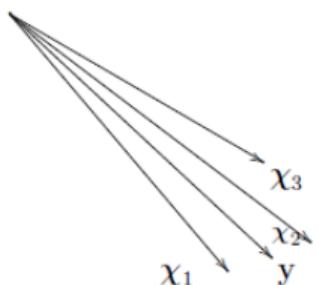
Multicollinear features and the forecast: possible configurations



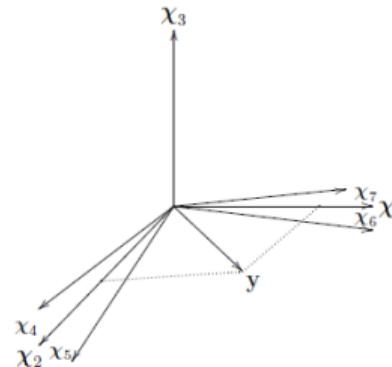
Inadequate and correlated



Adequate and random



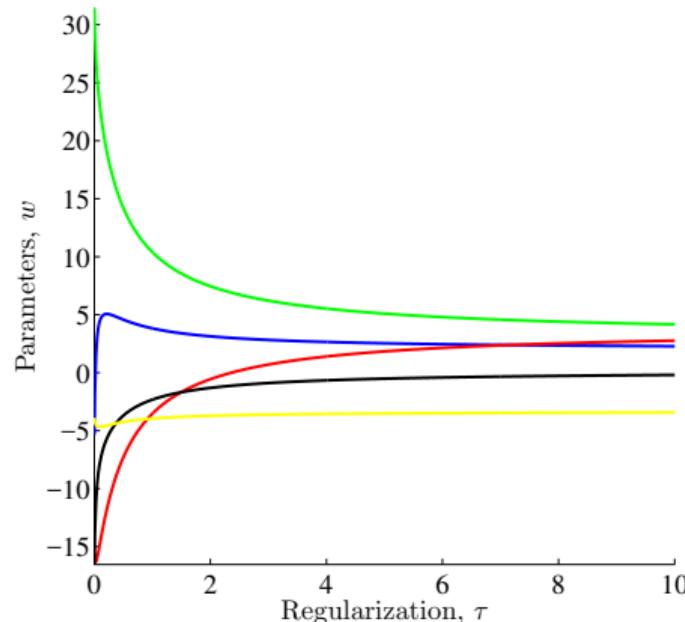
Adequate and redundant



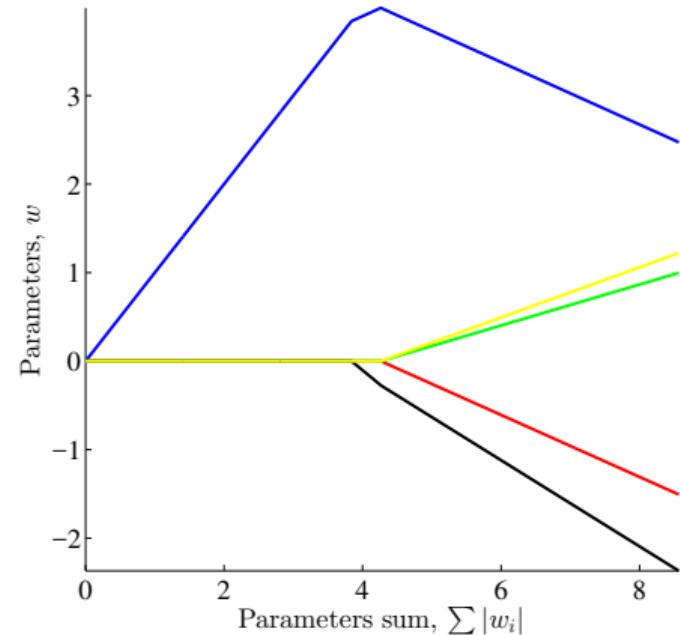
Adequate and correlated

Model parameter values with regularization

Vector-function $\mathbf{f} = \mathbf{f}(\mathbf{w}, \mathbf{X}) = [f(\mathbf{w}, \mathbf{x}_1), \dots, f(\mathbf{w}, \mathbf{x}_m)]^\top \in \mathbb{Y}^m$.



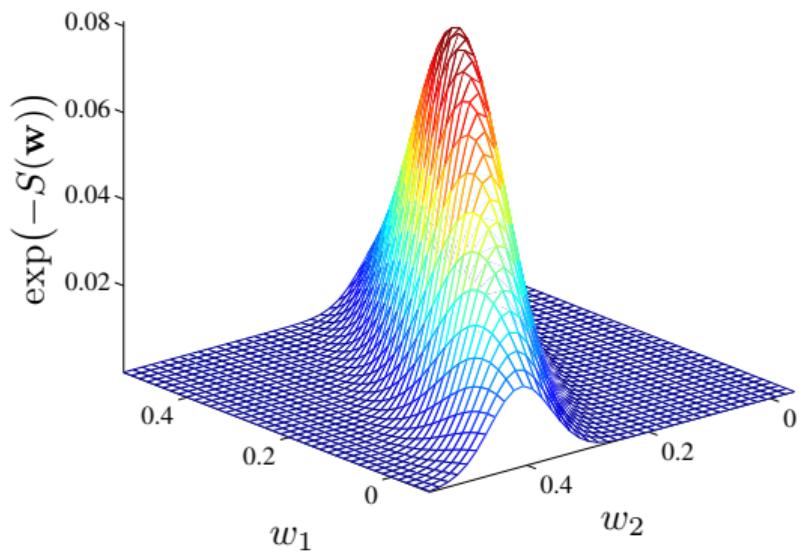
$$S(\mathbf{w}) = \|\mathbf{f}(\mathbf{w}, \mathbf{X}) - \mathbf{y}\|^2 + \gamma^2 \|\mathbf{w}\|^2$$



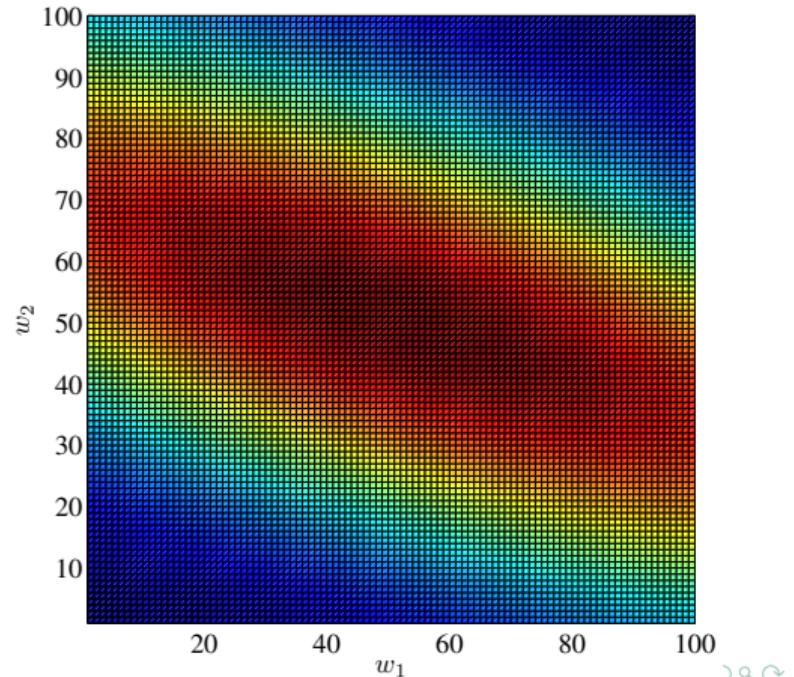
$$S(\mathbf{w}) = \|\mathbf{f}(\mathbf{w}, \mathbf{X}) - \mathbf{y}\|^2, T(\mathbf{w}) \leq \tau$$

Empirical distribution of model parameters

There given a sample $\{\mathbf{w}_1, \dots, \mathbf{w}_K\}$ of realizations of the m.r.v. \mathbf{w} and an error function $S(\mathbf{w}|\mathfrak{D}, \mathbf{f})$. Consider the set of points $\{s_k = \exp(-S(\mathbf{w}_k|\mathfrak{D}, \mathbf{f}))|k = 1, \dots, K\}$.



x- and y-axis: parameters \mathbf{w} , z-axis: $\exp(-S(\mathbf{w}))$.



Minimize number of similar and maximize number of relevant features

Introduce a feature selection method QP(Sim, Rel) to solve the optimization problem

$$\mathbf{a}^* = \arg \min_{\mathbf{a} \in \mathbb{B}^n} \mathbf{a}^\top \mathbf{Q} \mathbf{a} - \mathbf{b}^\top \mathbf{a},$$

where matrix $\mathbf{Q} \in \mathbb{R}^{n \times n}$ of pairwise similarities of features χ_i and χ_j is

$$\mathbf{Q} = [q_{ij}] = \text{Sim}(\chi_i, \chi_j) = \left| \frac{\text{Cov}(\chi_i, \chi_j)}{\sqrt{\text{Var}(\chi_i)\text{Var}(\chi_j)}} \right|$$

and vector $\mathbf{b} \in \mathbb{R}^n$ of feature relevances to the target is

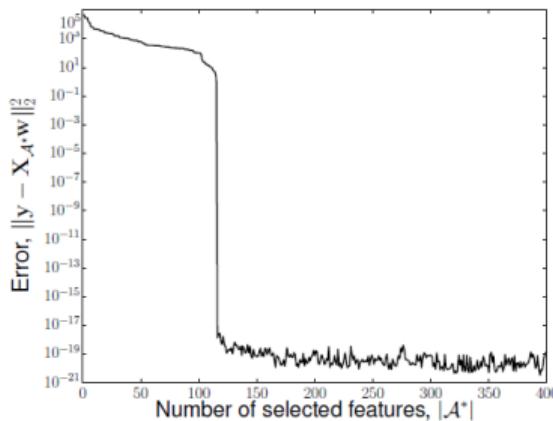
$$\mathbf{b} = [b_i] = \text{Rel}(\chi_i),$$

where elements b_i equal absolute values of the sample correlation coefficient between feature χ_i and the target vector \mathbf{y} .

Number of correlated features Sim \rightarrow min, number of correlated to the target Rel \rightarrow max.

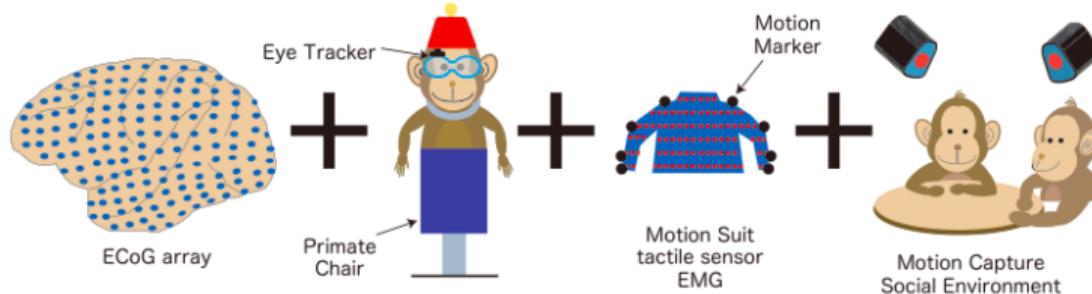
Evaluation criteria for the diesel NIR spectra data set

Method	C_p	RSS	$\ln \frac{\lambda_1}{\lambda_n}$ SVD	VIF	BIC
QP (ρ, ρ) ($\tau = 10^{-9}$)	-110	$1.37 \cdot 10^{-18}$	-25.7	$6.43 \cdot 10^6$	548.38
Genetic	-110.88	$7.68 \cdot 10^{-30}$	-24	$8.13 \cdot 10^5$	534.19
LARS	$3.22 \cdot 10^{21}$	$2.07 \cdot 10^{-7}$	-28.3	$7.94 \cdot 10^7$	529.47
Lasso	$2.5 \cdot 10^{28}$	1.61	-27.72	$1.03 \cdot 10^{21}$	1712.92
ElasticNet	$2.51 \cdot 10^{28}$	1.61	-27.72	$1.03 \cdot 10^{21}$	1712.92
Stepwise	$3.66 \cdot 10^{29}$	23.56	-36.78	$1.94 \cdot 10^{22}$	1919.23
Ridge	$1.59 \cdot 10^{28}$	1.02	-36.22	$1.07 \cdot 10^{22}$	$1.79 \cdot 10^3$



Dependence of residual norm on the number of selected features QP(Sim, Rel).

ECoG brain signals to forecast upper limbs' movements



- ▶ Neurotycho.org food-tracking dataset: 32 epidural electrodes captures brain signals of the monkey (ECoG),
- ▶ 11 sensors track movements of the hand, contralateral to the implant side,
- ▶ experiment duration is 15 minutes,
- ▶ the experimenter feeds the monkey ≈ 4.5 per minute,
- ▶ ECoG and motion data were sampled at 1KHz and 120Hz, respectively.

Model selection for time series forecasting

Thanks to the Chair of Intelligent Systems, MIPT

- ▶ Anastasia Motrenko
- ▶ Maria Popova
- ▶ Mikhail Kuznetsov
- ▶ Andrey Kulunchakov
- ▶ Alexandr Aduenko
- ▶ Mikhail Karasikov
- ▶ Arsentiy Kuzmin
- ▶ Radoslav Neychev
- ▶ Maria Stenina
- ▶ Alexey Goncharov
- ▶ Alexandr Katrutsa
- ▶ Roman Isachenko
- ▶ Oleg Bakhteev
- ▶ Maria Vladimirova

<http://machinelearning.ru>

strijov@phystech.edu